

```
In [1]: import torch
import numpy as np
import torch.nn as nn
from torchvision import transforms, utils, datasets
import matplotlib.pyplot as plt

transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,),(0.3081,)), # mean value = 0.1307, standard deviation value = 0.3081
])

data_path = './MNIST'

train_data = datasets.MNIST(root = data_path, train=False, download=True, transform= transform)
test_data = datasets.MNIST(root = data_path, train=True, download=True, transform= transform)

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz (http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz) to ./MNIST/MNIST/raw/train-images-idx3-ubyte.gz

HBox(children=(FloatProgress(value=1.0, bar_style='info', max=1.0), HTML(value='')))
```

Extracting ./MNIST/MNIST/raw/train-images-idx3-ubyte.gz to ./MNIST/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz (http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz) to ./MNIST/MNIST/raw/train-labels-idx1-ubyte.gz

HBox(children=(FloatProgress(value=1.0, bar\_style='info', max=1.0), HTML(value='')))

Extracting ./MNIST/MNIST/raw/train-labels-idx1-ubyte.gz to ./MNIST/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz (http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz) to ./MNIST/MNIST/raw/t10k-images-idx3-ubyte.gz

HBox(children=(FloatProgress(value=1.0, bar\_style='info', max=1.0), HTML(value='')))

Extracting ./MNIST/MNIST/raw/t10k-images-idx3-ubyte.gz to ./MNIST/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz (http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz) to ./MNIST/MNIST/raw/t10k-labels-idx1-ubyte.gz

HBox(children=(FloatProgress(value=1.0, bar\_style='info', max=1.0), HTML(value='')))

Extracting ./MNIST/MNIST/raw/t10k-labels-idx1-ubyte.gz to ./MNIST/MNIST/raw

Processing...

Done!

/usr/local/lib/python3.6/dist-packages/torchvision/datasets/mnist.py:480: UserWarning: The given NumPy array is not writeable, and PyTorch does not support non-writeable tensors. This means you can write to the underlying (supposedly non-writeable) NumPy array using the tensor. You may want to copy the array to protect its data or make it writeable before converting it to a tensor. This type of warning will be suppressed for the rest of this program. (Triggered internally at /pytorch/torch/csrc/utils/tensor\_numpy.cpp:141.)

```
    return torch.from_numpy(parsed.astype(m[2], copy=False)).view(*s)

cuda:0
```

```
In [2]: print("the number of your training data (must be 10,000) = ", train_data.__len__())
print("hte number of your testing data (must be 60,000) = ", test_data.__len__())

the number of your training data (must be 10,000) = 10000
hte number of your testing data (must be 60,000) = 60000
```

```
In [3]: class Classification(nn.Module):
def __init__(self, param1, param2, dr_rate=0):
    super(Classification, self).__init__()
    self.param1 = param1
    self.param2 = param2
    # construct layers for a neural network
    self.classifier1 = nn.Sequential(
        nn.Linear(in_features=28*28, out_features=self.param1),
        nn.ReLU(),
        nn.Dropout(p=dr_rate),
    )
    self.classifier2 = nn.Sequential(
        nn.Linear(in_features=self.param1, out_features=self.param2),
        nn.ReLU(),
        nn.Dropout(p=dr_rate),
    )
    self.classifier3 = nn.Sequential(
        nn.Linear(in_features=self.param2, out_features=10),
        nn.LogSoftmax(dim=1),
    )

def forward(self, inputs):
    x = inputs.view(inputs.size(0), -1) # [batchSize, 28*28]
    x = self.classifier1(x) # [batchSize, 20*20]
    x = self.classifier2(x) # [batchSize, 10*10]
    out = self.classifier3(x) # [batchSize, 10]

    return out
```

```
In [4]: lr=0.00001

classification = Classification(10,10,0).to(device)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(classification.parameters(), lr=lr)
```

```
In [5]: def run_epoch (train_data, test_data):

    tr_loss = 0
    tr_acc = 0
    iter = len(train_data)
    classification.train()
    for img_i, label_i in train_data: #{
        img_i, label_i = img_i.to(device), label_i.to(device)

        optimizer.zero_grad()
        # Forward
        y_pred = classification.forward(img_i.view(-1, 28*28))
        ps = torch.exp(y_pred)
        correct = (label_i == ps.max(dim=1)[1])
        # Loss computation
        loss = criterion(y_pred, label_i)
        # Backward
        loss.backward()
        # Optimize for img_i
        optimizer.step()
        tr_loss += loss.item()
        tr_acc += correct.type(torch.FloatTensor).mean()

    #}
    tr_loss /= iter
    tr_acc /= iter

    test_loss = 0
    test_acc = 0
    iter_test = len(test_data)

    classification.eval()
    for img_j, label_j in test_data:
        img_j, label_j = img_j.to(device), label_j.to(device)
        correct = 0
        with torch.autograd.no_grad():
            predicted = classification.forward(img_j.view(-1, 28*28))
            ps = torch.exp(predicted)
            correct = (label_j == ps.max(dim=1)[1])
            test_loss += criterion(predicted, label_j).item()
            test_acc += correct.type(torch.FloatTensor).mean()

    test_loss /= iter_test
    test_acc /= iter_test
    return tr_loss, tr_acc, test_loss, test_acc

final_train_loss = []
final_test_loss = []
final_train_acc = []
final_test_acc = []

def run(batch_size, epochs, param1, param2, dr_rate): #{
    global optimizer, criterion, classification
    classification = Classification(param1, param2, dr_rate).to(device)
    criterion = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(classification.parameters(), lr=lr)

    global final_train_loss, final_test_loss, final_train_acc, final_test_acc
    train_data_loader = torch.utils.data.DataLoader(train_data, batch_size, shuffle=False)
    test_data_loader = torch.utils.data.DataLoader(test_data, batch_size, shuffle=False)
    mini_batch_data, mini_batch_label = next(iter(train_data_loader))

    train_loss = []
    test_loss = []
    train_acc = []
    test_acc = []

    for epoch in range(epochs): #{
        tr_loss, tr_acc, te_loss, te_acc = run_epoch(train_data_loader, test_data_loader)
        train_loss.append(tr_loss)
        train_acc.append(tr_acc)
        test_loss.append(te_loss)
        test_acc.append(te_acc)
    #}

    final_train_loss.append(train_loss[-1])
    final_test_loss.append(test_loss[-1])
    final_train_acc.append(train_acc[-1])
    final_test_acc.append(test_acc[-1])


    return train_loss, test_loss, train_acc, test_acc

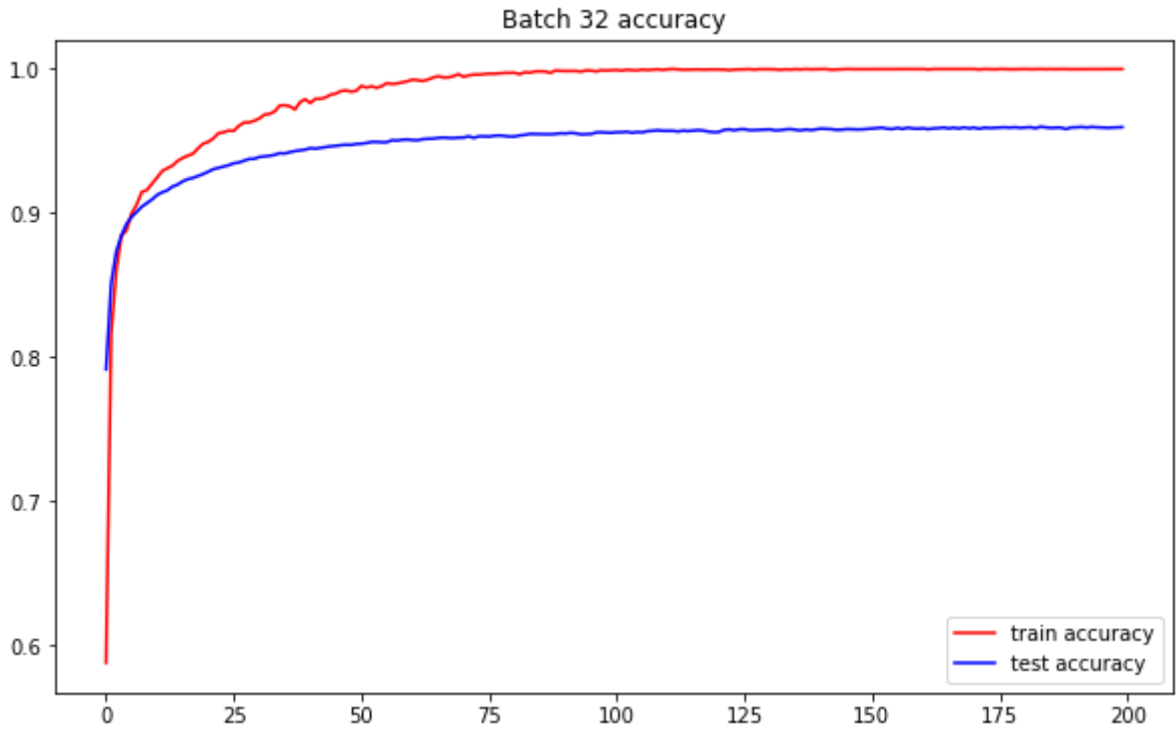
#}

train_loss_32, test_loss_32, train_acc_32, test_acc_32 = run(32, 200,2048,1024, 0.4)

# Plot image
plt.figure(0, figsize=(10,6))
plt.plot(train_loss_32, label='train loss', c='r')
plt.plot(test_loss_32, label='test loss', c='b')
plt.title(f'Batch 32 loss')
plt.legend()
plt.show()

plt.figure(1, figsize=(10,6))
plt.plot(train_acc_32, label='train accuracy', c='r')
plt.plot(test_acc_32, label='test accuracy', c='b')
plt.title(f'Batch 32 accuracy')
plt.legend()
plt.show()
```

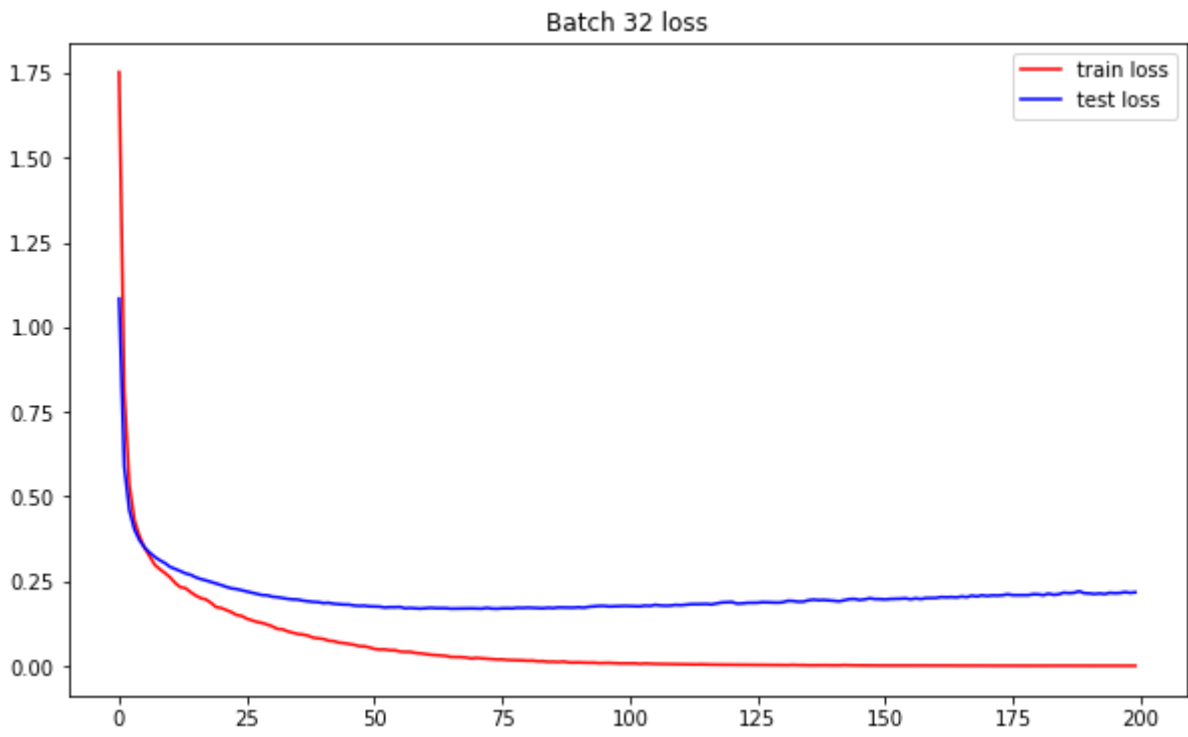




## Output using the dataset

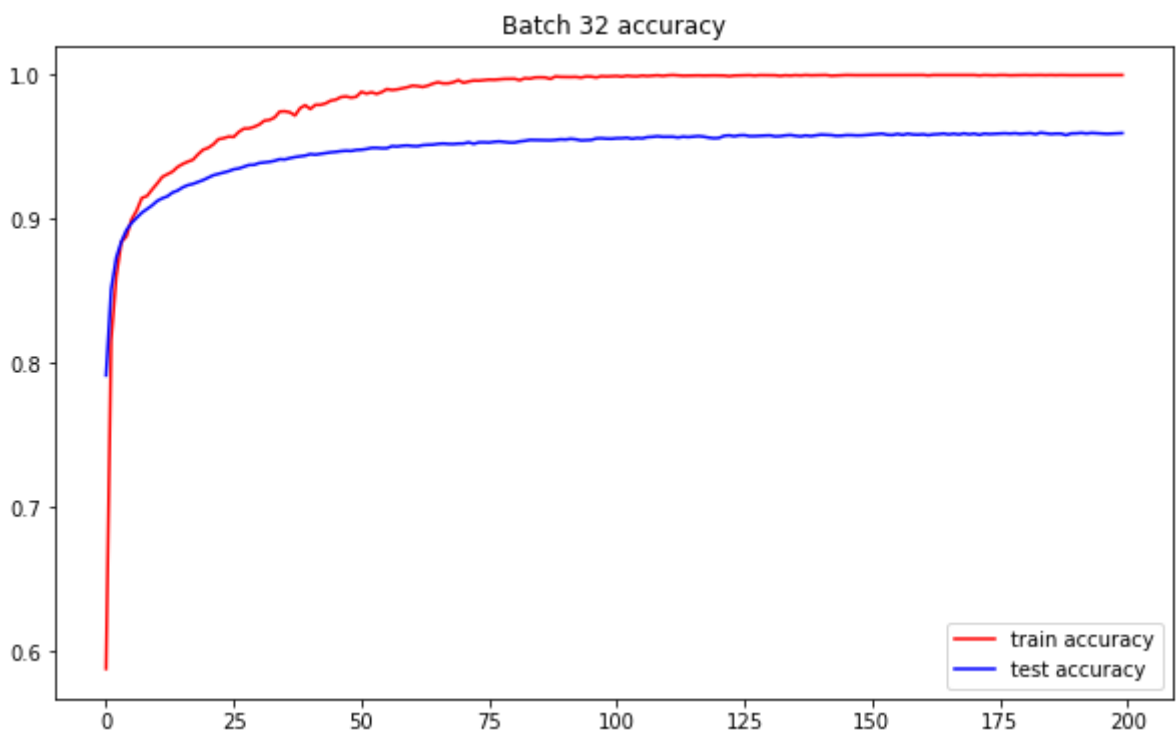
1. Plot the training and testing losses over epochs

```
In [6]: plt.figure(0, figsize=(10,6))
plt.plot(train_loss_32, label='train loss', c='r')
plt.plot(test_loss_32, label='test loss', c='b')
plt.title(f'Batch 32 loss')
plt.legend()
plt.show()
```



2. Plot the training and testing accuracies over epochs

```
In [7]: plt.figure(1, figsize=(10,6))
plt.plot(train_acc_32, label='train accuracy', c='r')
plt.plot(test_acc_32, label='test accuracy', c='b')
plt.title(f'Batch 32 accuracy')
plt.legend()
plt.show()
```



3. Print the final training and testing losses at convergence

```
In [8]: print('Loss')
print(f'training: {train_loss_32[-1]}')
print(f'testing: {test_loss_32[-1]}')
```

Loss  
training: 0.0006691269123527549  
testing: 0.21748759474004856

4. Print the final training and testing accuracies at convergence

```
In [9]: print('Accuracy')
print(f'training: {train_acc_32[-1]}')
print(f'testing: {test_acc_32[-1]}')
```

Accuracy  
training: 1.0  
testing: 0.9596999883651733