

```
In [1]: import torch
import numpy as np
import torch.nn as nn
from torchvision import transforms, utils, datasets
import matplotlib.pyplot as plt

transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,)), # mean value = 0.1307, standard deviation value = 0.3081
])

data_path = './MNIST'

train_data = datasets.MNIST(root = data_path, train=False, download=True, transform= transform)
test_data = datasets.MNIST(root = data_path, train=True, download=True, transform= transform)

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz (http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz) to ./MNIST/MNIST/raw/train-images-idx3-ubyte.gz

HBox(children=(FloatProgress(value=1.0, bar_style='info', max=1.0), HTML(value='')))

Extracting ./MNIST/MNIST/raw/train-images-idx3-ubyte.gz to ./MNIST/MNIST/raw
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz (http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz) to ./MNIST/MNIST/raw/train-labels-idx1-ubyte.gz

HBox(children=(FloatProgress(value=1.0, bar_style='info', max=1.0), HTML(value='')))

Extracting ./MNIST/MNIST/raw/train-labels-idx1-ubyte.gz to ./MNIST/MNIST/raw
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz (http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz) to ./MNIST/MNIST/raw/t10k-images-idx3-ubyte.gz

HBox(children=(FloatProgress(value=1.0, bar_style='info', max=1.0), HTML(value='')))

Extracting ./MNIST/MNIST/raw/t10k-images-idx3-ubyte.gz to ./MNIST/MNIST/raw
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz (http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz) to ./MNIST/MNIST/raw/t10k-labels-idx1-ubyte.gz

HBox(children=(FloatProgress(value=1.0, bar_style='info', max=1.0), HTML(value='')))

Extracting ./MNIST/MNIST/raw/t10k-labels-idx1-ubyte.gz to ./MNIST/MNIST/raw
Processing...
Done!

/usr/local/lib/python3.6/dist-packages/torchvision/datasets/mnist.py:480: UserWarning: The given NumPy array is not writeable, and PyTorch does not support non-writeable tensors. This means you can write to the underlying (supposedly non-writeable) NumPy array using the tensor. You may want to copy the array to protect its data or make it writeable before converting it to a tensor. This type of warning will be suppressed for the rest of this program. (Triggered internally at /pytorch/torch/csrc/utils/tensor_numpy.cpp:141.)
    return torch.from_numpy(parsed.astype(m[2], copy=False)).view(*s)

cuda:0
```

```
In [2]: print("the number of your training data (must be 10,000) = ", train_data.__len__())
print("hte number of your testing data (must be 60,000) = ", test_data.__len__())

the number of your training data (must be 10,000) = 10000
hte number of your testing data (must be 60,000) = 60000
```

```
In [3]: class CNN(nn.Module):
    def __init__(self,dr_rate=0):
        super(CNN,self).__init__()
        self.layer = nn.Sequential(
            nn.Conv2d(1, 32, 3, padding=1),
            nn.ReLU(),
            nn.Dropout(p=dr_rate),
            nn.Conv2d(32, 64, 3,padding=1),
            nn.ReLU(),
            nn.Dropout(p=dr_rate),
            nn.Conv2d(64, 128, 3,padding=1),
            nn.ReLU(),
            nn.Dropout(p=dr_rate),
            nn.MaxPool2d(2, 2),
            nn.Conv2d(128, 256, 3,padding=1),
            nn.ReLU(),
            nn.Dropout(p=dr_rate),
            nn.Conv2d(256, 256, 3,padding=1),
            nn.ReLU(),
            nn.Dropout(p=dr_rate),
            nn.MaxPool2d(2, 2),
        )
        self.fc_layer = nn.Sequential(
            nn.Linear(256*7*7, 128),
            nn.ReLU(),
            nn.Dropout(p=dr_rate),
            nn.Linear(128, 64),
            nn.ReLU(),
            nn.Linear(64, 10),
            nn.LogSoftmax(dim=1),
        )

    def forward(self,x):
        out = self.layer(x)
        # print('a')
        out = out.view(50, -1)
        # print('b')
        out = self.fc_layer(out)
        # print('c')
        return out
```

```
In [4]: lr=0.001
epochs=100
dr_rate = 0.5

classification = CNN(dr_rate).to(device)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(classification.parameters(), lr=lr)
```

```
In [5]: def run_epoch (train_data, test_data):

    tr_loss = 0
    tr_acc = 0
    iter = len(train_data)
    classification.train()
    for img_i, label_i in train_data: #{
        img_i, label_i = img_i.to(device), label_i.to(device)

        optimizer.zero_grad()
        # Forward
        y_pred = classification.forward(img_i)
        correct = (label_i == y_pred.max(dim=1)[1])
        # Loss computation
        loss = criterion(y_pred, label_i)
        # Backward
        loss.backward()
        # Optimize for img_i
        optimizer.step()
        tr_loss += loss.item()
        tr_acc += correct.type(torch.FloatTensor).mean()
    #}
    tr_loss /= iter
    tr_acc /= iter
    test_loss = 0
    test_acc = 0
    iter_test = len(test_data)

    classification.eval()
    correct = 0
    total = 0
    with torch.autograd.no_grad():
        for img_j, label_j in test_data:
            img_j, label_j = img_j.to(device), label_j.to(device)
            predicted = classification.forward(img_j)
            _, output_idx = torch.max(predicted, 1)
            total += label_j.size(0)
            correct += (output_idx == label_j).sum().float()
            test_loss += criterion(predicted, label_j).item()

    test_loss /= iter_test
    test_acc = correct / total
    tr_acc = tr_acc.item()
    test_acc = test_acc.item()
    return tr_loss, tr_acc, test_loss, test_acc

def run(batch_size, epochs): #{


    train_data_loader = torch.utils.data.DataLoader(train_data, batch_size, shuffle=True)
    test_data_loader = torch.utils.data.DataLoader(test_data, batch_size, shuffle=False)

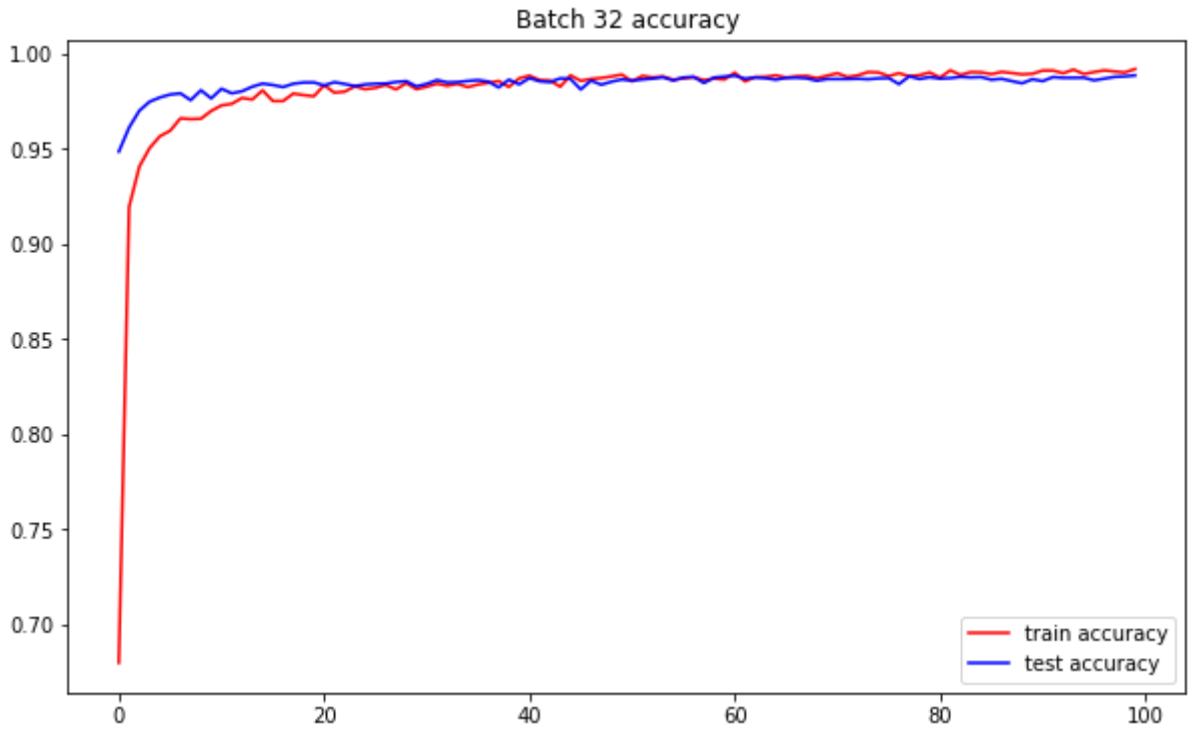
    train_loss = []
    test_loss = []
    train_acc = []
    test_acc = []
    for epoch in range(epochs): #{
        tr_loss, tr_acc, te_loss, te_acc = run_epoch(train_data_loader, test_data_loader)
        train_loss.append(tr_loss)
        train_acc.append(tr_acc)
        test_loss.append(te_loss)
        test_acc.append(te_acc)
    #}
    return train_loss, test_loss, train_acc, test_acc

#}
train_loss_32, test_loss_32, train_acc_32, test_acc_32 = run(50, epochs)

# Plot image
plt.figure(0, figsize=(10,6))
plt.plot(train_loss_32, label='train loss', c='r')
plt.plot(test_loss_32, label='test loss', c='b')
plt.title(f'Batch 32 loss')
plt.legend()
plt.show()

plt.figure(1, figsize=(10,6))
plt.plot(train_acc_32, label='train accuracy', c='r')
plt.plot(test_acc_32, label='test accuracy', c='b')
plt.title(f'Batch 32 accuracy')
plt.legend()
plt.show()
```

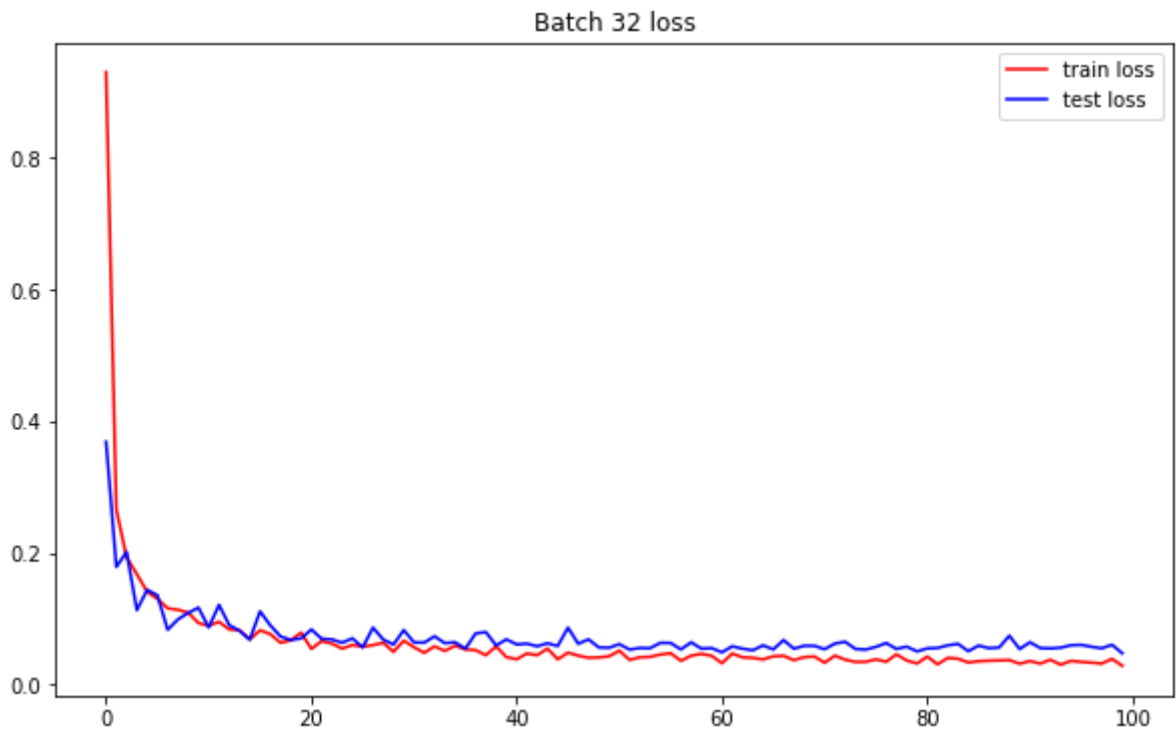




## Output using the dataset


1. Plot the training and testing losses over epochs

```
In [6]: plt.figure(0, figsize=(10,6))
plt.plot(train_loss_32, label='train loss', c='r')
plt.plot(test_loss_32, label='test loss', c='b')
plt.title(f'Batch 32 loss')
plt.legend()
plt.show()
```



2. Plot the training and testing accuracies over epochs

```
In [7]: plt.figure(1, figsize=(10,6))
plt.plot(train_acc_32, label='train accuracy', c='r')
plt.plot(test_acc_32, label='test accuracy', c='b')
plt.title(f'Batch 32 accuracy')
plt.legend()
plt.show()
```



3. Print the final training and testing losses at convergence

```
In [8]: print('Loss')
print(f'training: {round(train_loss_32[-1], 5)}')
print(f'testing: {round(test_loss_32[-1], 5)}')
```

```
Loss
training: 0.02891
testing: 0.048
```

4. Print the final training and testing accuracies at convergence

```
In [9]: print('Accuracy')
print(f'training: {round(train_acc_32[-1], 5)}')
print(f'testing: {round(test_acc_32[-1], 5)}')
```

```
Accuracy
training: 0.9917
testing: 0.98835
```

5. Print the testing accuracies within the last 10 epochs

```
In [10]: for i in range(10): #{
        print(f'epoch = {epochs-9+i}] {round(test_acc_32[-10+i], 5)}')
        #}
```

```
[epoch = 91] 0.98535
[epoch = 92] 0.98747
[epoch = 93] 0.98708
[epoch = 94] 0.98708
[epoch = 95] 0.98715
[epoch = 96] 0.98578
[epoch = 97] 0.98665
[epoch = 98] 0.98758
[epoch = 99] 0.98785
[epoch = 100] 0.98835
```

```
In [10]:
```